

Data Science and Visualization (DSV, F23)

9. Association Rules

Hua Lu

<https://luhua.ruc.dk>; luhua@ruc.dk

PLIS, IMT, RUC

Recommendation in Amazon

- Two types of recommendation
 - 90% buyers who bought A also bought B.
 - Since you've bought A, you may also want B.

Frequently Bought Together

Frequent itemset



Price for all three: \$67.41

Add all three to Cart

Add all three to Wish List

Show availability and shipping details

- ☒ This item: Big Data: A Revolution That Will Transform How We Live, Work, and Think by Viktor Mayer-Schönberger Paperback \$10.61
- ☒ Data Science for Business: What you need to know about data mining and data-analytic thinking by Foster Provost Paperback \$37.99
- ☒ Predictive Analytics: The Power to Predict Who Will Click, Buy, Lie, or Die by Eric Siegel Hardcover \$18.81

Add to Wish List

Have one to sell?


Sell on Amazon

Customers Who Bought This Item Also Bought

Association rules



Data Science for Business:
What you need...
› Foster Provost
★★★★★ 103
#1 Best Seller in Database
Storage & Design
Paperback
\$37.99 ✓Prime



Predictive Analytics: The
Power to Predict Who...
› Eric Siegel
★★★★☆ 235
Hardcover
\$18.81 ✓Prime



Big Data: Understanding
How Data Powers Big...
› Bill Schmarzo
★★★★★ 5
Paperback
\$28.94 ✓Prime



Big Data For Dummies
› Judith Hurwitz
★★★★☆ 31
Paperback
\$19.81 ✓Prime



Digital Sociology
Deborah Lupton
Paperback
\$46.95 ✓Prime



Big Data: Using SMART
Big Data, Analytics...
› Bernard Marr
★★★★☆ 9
Paperback
\$16.74 ✓Prime

Page

Agenda

- Problem definition
 - Support, confidence, lift, and association rule
 - Frequent itemsets
 - Steps for association rule mining
- Apriori principle, Apriori algorithm
- Deriving association rules from frequent itemsets

Market Basket Data

- Large set of *items*, i.e., things sold in a supermarket
- Large set of *baskets*, each a small subset of items, i.e., things that one customer buys in one **transaction**
- Transaction table **T**: market-basket data
 - Each record is a transaction, containing a set of items
 - Many-to-many mapping (association) between items and baskets
- What can we do with this type of data?
 - E.g., counting whether the combination {Milk, Bread} is *frequent* or not

TID	Items
1	{Milk, Bread, Beer, Diapers}
2	{Bread, Eggs}
3	{Bread, Diapers}
4	{Milk, Bread, Cola}
5	{Milk, Bread, Diapers}

Transaction table

What Is Association Rule Mining?

- Finding **frequent patterns** and **associations** (rules) among sets of items in a transaction table
- Motivation (market basket analysis):
 - How likely is that the customers buying *milk* are also buying *bread*?
 - Such rules help retailers making decisions
 - Plan the shelf space: placing milk close to bread, more convenient for the customers
 - Offer promotions/discounts for those products together

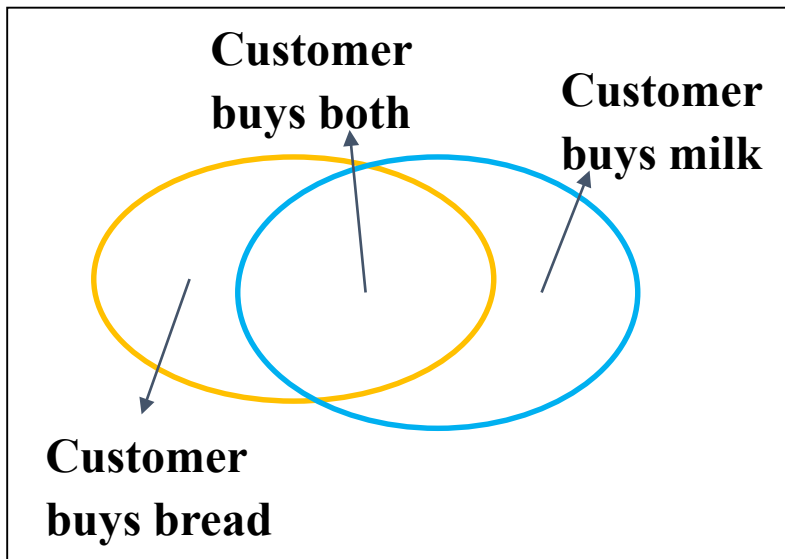
What Is an Association Rule?

- An **association rule** correlates (associates) the presence of one set of items with that of another set of items
- Examples
 - Rule form: $\text{Body} \Rightarrow \text{Head} [\text{support}, \text{confidence}]$
 - $\text{milk} \Rightarrow \text{bread} [5\%, 70\%]$
 - 5% of transactions buy both milk and bread
 - transactions that buy milk have 70% chance of buying also bread
- Applications: basket data analysis, catalog design
 - $* \Rightarrow \text{chocolate}$ (How to boost the sales of chocolate)
 - $\text{Home Electronics} \Rightarrow *$ (What other products should the store stock up?)

Interesting Rules

- A rule is said to be **interesting** (or valid) when:
 - Its items appear frequently in the transaction table (**support**)
 - It holds with a high probability (**confidence**)

Example: $milk \Rightarrow bread$



NB: X and Y are itemsets.

Find all the rules $X \Rightarrow Y$ with confidence and support above given thresholds

- **support** s , probability that a transaction contains $X \cup Y$
- **confidence** c , conditional probability that a transaction having X also contains Y

Example (1)

- Find the **support** and **confidence** of the rule: $\{B,D\} \Rightarrow \{A\}$
- Support value of $sup(ABD)$:
 - percentage of tuples with $\{A,B,D\}$
 $= (3/4) * 100\% = 75\%$
- Confidence value of $conf(BD \Rightarrow A)$
$$\frac{\text{number of transactions that contain } \{A, B, D\}}{\text{number of transactions that contain } \{B, D\}} = \frac{3}{3} = 100\%$$

TID	items bought
100	{F,A,D,B}
200	{D,A,C,E,B}
300	{C,A,B,E}
400	{B,A,D}

$$\text{prob}(Y | X) = \frac{\text{prob}(X \cup Y)}{\text{prob}(X)}$$
$$\text{conf}(X \Rightarrow Y) = \frac{\text{sup}(X \cup Y)}{\text{sup}(X)} = \frac{\text{frequency}(X \cup Y)}{\text{frequency}(X)}$$

Example (2)

- Find interesting rules

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

Thresholds:

Min. support 50%
Min. confidence 50%

Frequent Itemset	Support
{A}	75%
{B}	50%
{C}	50%
{A,C}	50%

“ $A \Rightarrow C$ ” is a valid rule because:

support = support($\{A \cup C\}$) = $2/4 = 50\%$

confidence = support($\{A \cup C\}$)/support($\{A\}$) = $50\%/75\% = 66.6\%$

Lift of A Rule

- **Lift**($X \Rightarrow Y$) = $\text{confidence}(X \Rightarrow Y) / \text{support}(Y)$
= $\text{support}(X \cup Y) / (\text{support}(X) * \text{support}(Y))$
= $(\text{frequency}(X \cup Y) * |T|) / (\text{frequency}(X) * \text{frequency}(Y))$
- Lift($X \Rightarrow Y$) refers to **the increase in the ratio of sale of Y when X is sold**
 - **Lift = 1**: No association between products X and Y.
 - **Lift > 1**: Products X and Y are more likely to be bought together.
 - **Lift < 1**: The two products are unlikely to be bought together.

Example of Lift

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

Frequent Itemset	Support
{A}	75%
{B}	50%
{C}	50%
{A,C}	50%

Rule $A \Rightarrow C$:

support = $\text{support}(\{A \cup C\}) = 2/4 = 50\%$

confidence = $\text{support}(\{A \cup C\}) / \text{support}(\{A\}) = 50\% / 75\% = 66.6\%$

lift = $\text{confidence}(A \Rightarrow C) / \text{support}(C) = 66.6\% / 50\% = 1.333$




Lift's meaning: the likelihood of buying a A and C *together* is 1.33 times more than the likelihood of just buying the C.

A Real Application of Association Rules

- Amazon's recommendation
 - 90% buyers who bought A also bought B.
 - Since you've bought A, you may also want B.

Make recommendations based on rules of high support, confidence and lift.

Frequently Bought Together

 +  + 

Price for all three: **\$67.41**

[Add all three to Cart](#)

[Add all three to Wish List](#)

[Show availability and shipping details](#)


[Add to Wish List](#)

Have one to sell? [Sell on Amazon](#)


- ☒ **This item:** Big Data: A Revolution That Will Transform How We Live, Work, and Think by Viktor Mayer-Schönberger Paperback **\$10.61**
- ☒ Data Science for Business: What you need to know about data mining and data-analytic thinking by Foster Provost Paperback **\$37.99**
- ☒ Predictive Analytics: The Power to Predict Who Will Click, Buy, Lie, or Die by Eric Siegel Hardcover **\$18.81**

Customers Who Bought This Item Also Bought


Page 1




Data Science for Business: What you need...
› Foster Provost
★★★★★ 103
#1 Best Seller in Database Storage & Design
Paperback
\$37.99 ✓Prime




Predictive Analytics: The Power to Predict Who...
› Eric Siegel
★★★★☆ 235
Hardcover
\$18.81 ✓Prime




Big Data: Understanding How Data Powers Big...
› Bill Schmarzo
★★★★★ 5
Paperback
\$28.94 ✓Prime



Big Data For Dummies
› Judith Hurwitz
★★★★☆ 31
Paperback
\$19.81 ✓Prime



Digital Sociology
Deborah Lupton
Paperback
\$46.95 ✓Prime



Big Data: Using SMART Big Data, Analytics...
› Bernard Marr
★★★★☆ 9
Paperback
\$16.74 ✓Prime

Causality vs. Correlation

- Causality
 - From the very first day, humans are curious about *why*.
 - With big data, it may be very hard to see the exact reasons.
- Correlation
 - Instead, we can find interesting patterns or associations of different things from big data.
 - Probability instead of certainty (not totally random).
 - Association rule mining.
- **NB:** Association rules are empiricism! What they tell may not be the true cause and effect.

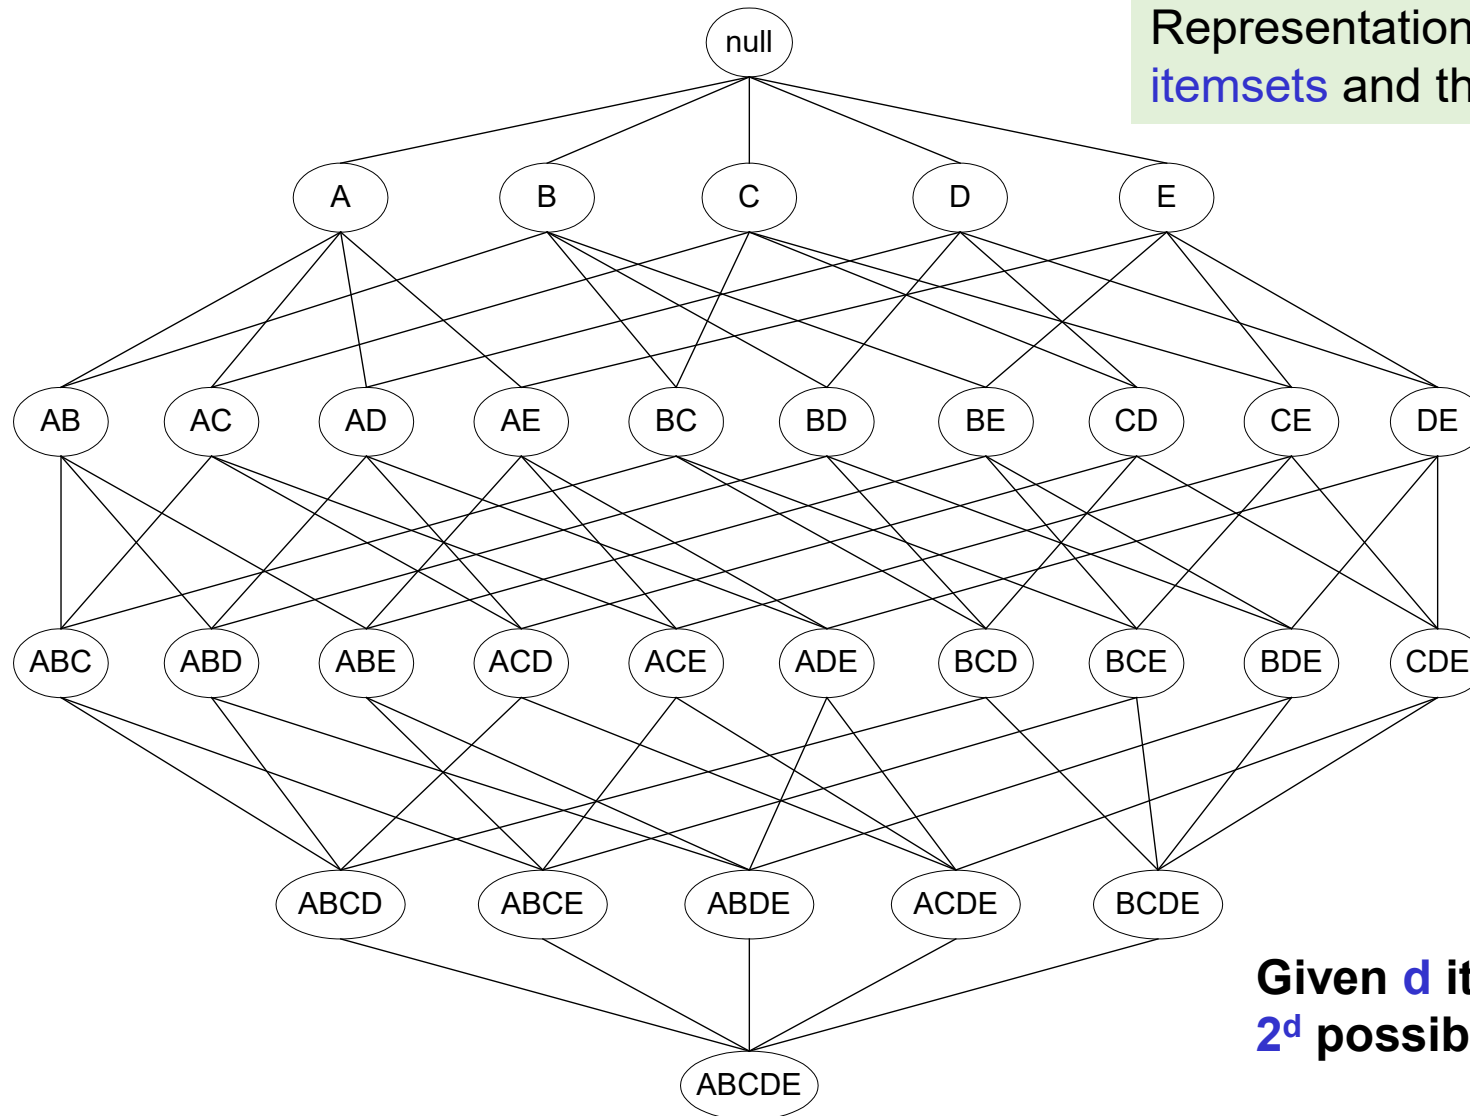
Steps of Association Rule Mining

1. Find the *frequent itemsets*
 - The sets of items that have minimum support
 - How to do this efficiently?
2. Use the frequent itemsets to generate association rules

Mining Frequent Itemsets

- **Input:** A set of transactions T , over a set of items I
- **Output:** All itemsets with items in I having
 - support \geq minsup (support threshold)
- Problem parameters:
 - $N = |T|$: number of transactions
 - $d = |I|$: number of (distinct) items
 - w : max width of a transaction
 - Number of possible itemsets: $M = 2^d$
- Scale of the problem:
 - WalMart sells 100,000 items and can store billions of baskets.
 - The Web has billions of words and many billions of pages.

The Itemset Lattice



Representation of all possible
itemsets and their relationships

Given **d** items, there are
 2^d possible itemsets

Agenda

- Problem definition
- Apriori principle, Apriori algorithm
- Deriving association rules from frequent itemsets

The Apriori Principle

- Main observations: $\forall X, Y: X \subseteq Y \Rightarrow s(X) \geq s(Y)$
 - If an itemset is frequent, so are its subsets
 - If an itemset is infrequent, so are its supersets
- The **Apriori** principle: *A subset of a frequent itemset must also be a frequent itemset*
 - E.g., if $\{AB\}$ is a frequent itemset, both $\{A\}$ and $\{B\}$ must be a frequent itemset
 - Iteratively find frequent itemsets with cardinality from 1 to m (m -itemset):
Use frequent k -itemsets to explore $(k+1)$ -itemsets

Illustration of Apriori Principle

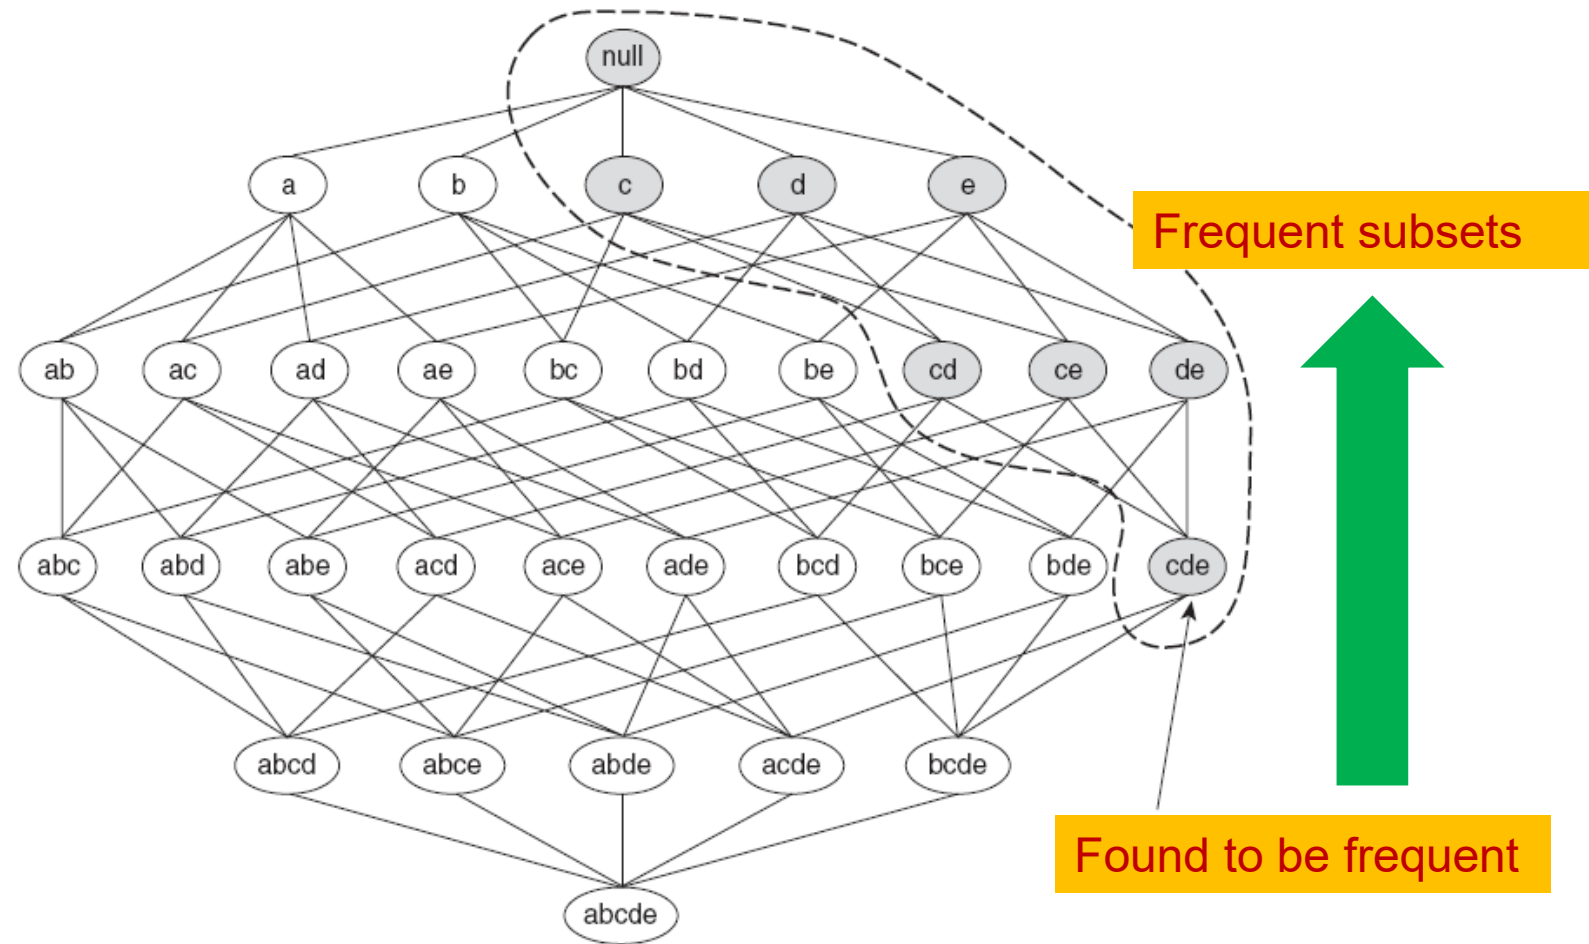
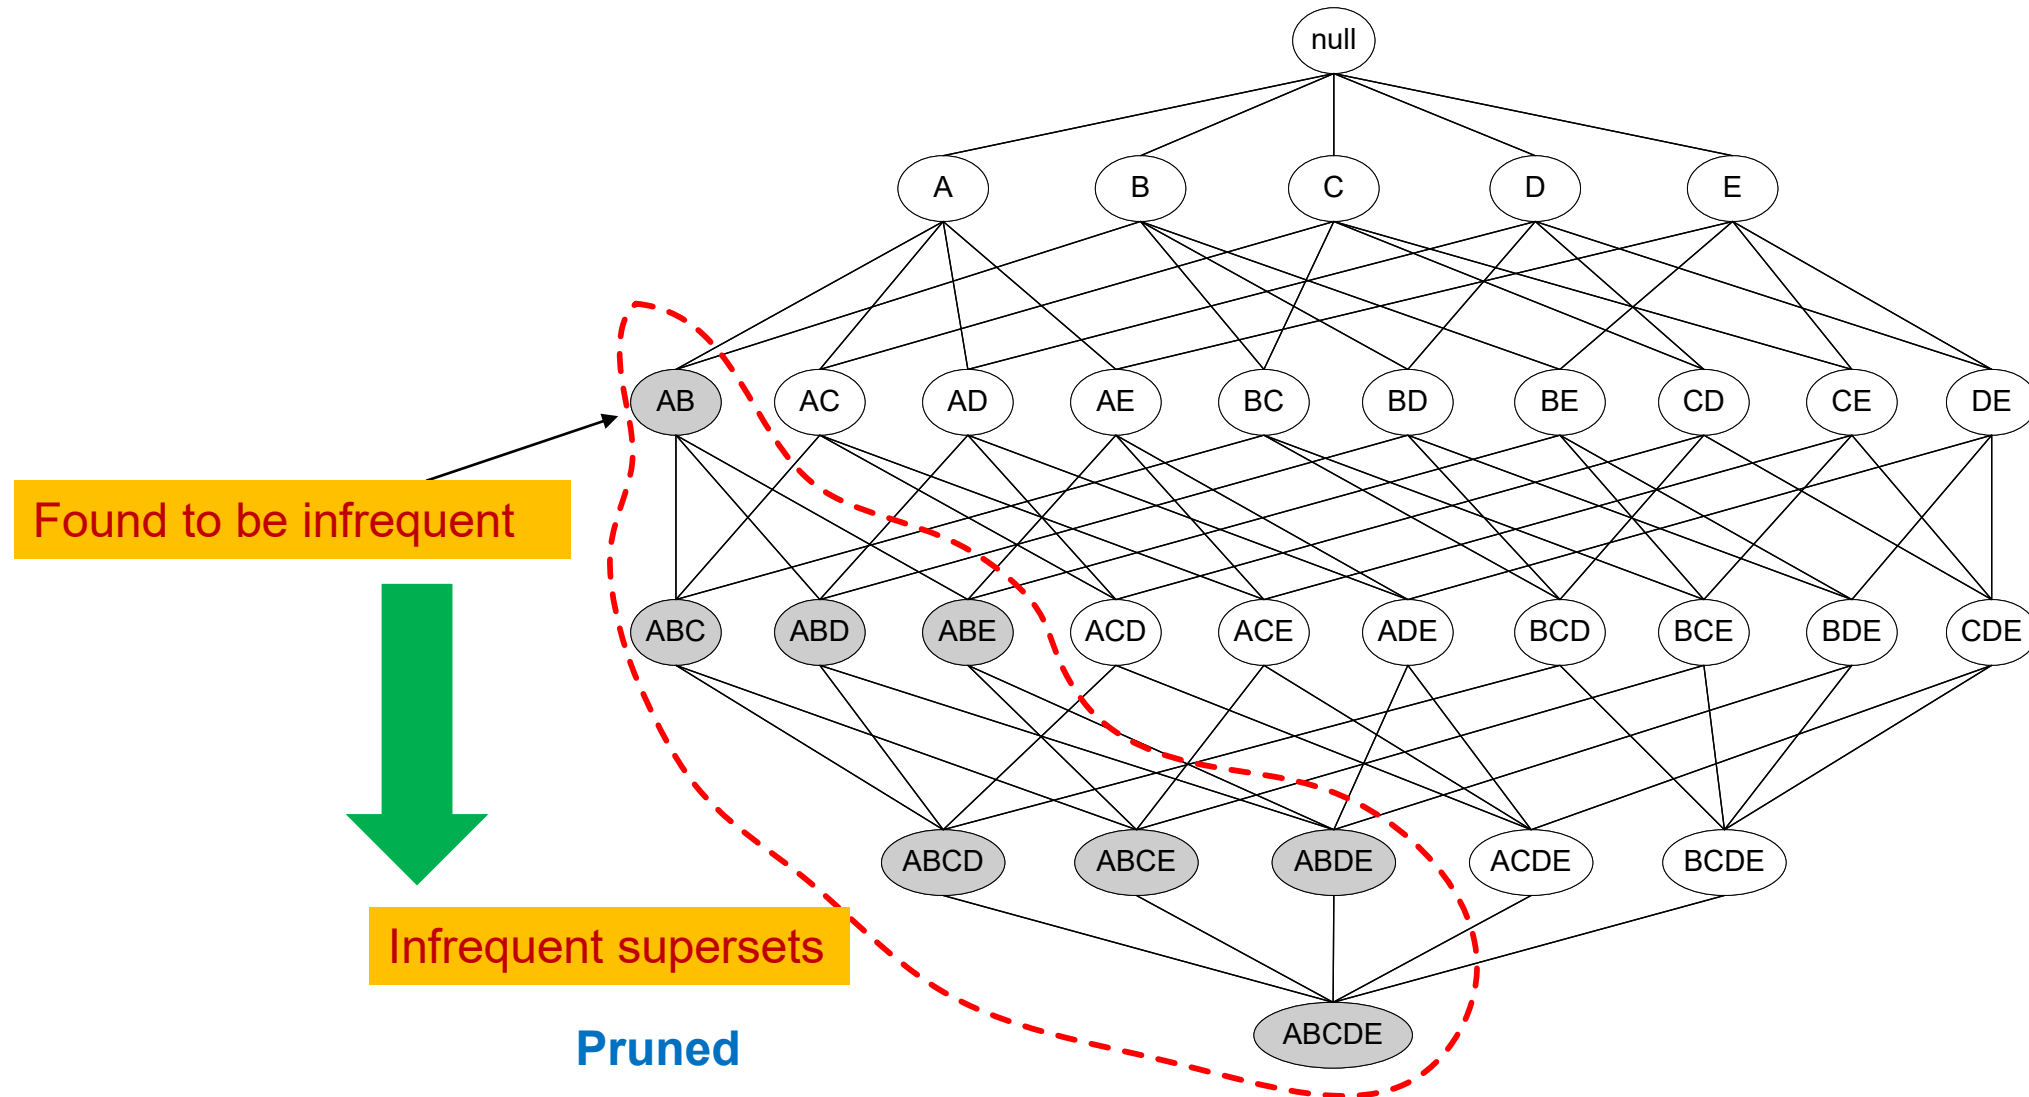


Figure 6.3. An illustration of the *Apriori* principle. If $\{c, d, e\}$ is frequent, then all subsets of this itemset are frequent.

Illustration of Apriori Principle (cont.)



Level-wise Process of Apriori Principle

Level 4 (frequent quadruples): $\{....\}$

Level 3 (frequent triplets): $\{ABD\}, \{BDF\}$

Level 2 (frequent pairs): $\{AB\}, \{AD\}, \{BD\}, \{BF\}, \{DF\}$

Level 1 (frequent items): $\{A\}, \{B\}, \{D\}, \{F\}$

Remember:

All subsets of a frequent itemset must be frequent as well

Question: Can ADF be frequent?

NO: because AF is not frequent

The Apriori Algorithm

Advanced

- Notations
 - C_k : Candidate itemset of size k
 - L_k : Frequent itemset of size k
- Important steps in candidate generation
 - **Prune Step**: Any k -itemset that is not frequent cannot be a subset of a frequent $(k+1)$ -itemset
 - **Join Step**: C_{k+1} is generated by joining L_k with itself

```
 $C_I = \{\{\text{item}_1\}, \dots, \{\text{item}_N\}\};$   
for ( $k = 1$ ;  $L_k \neq \emptyset$ ;  $k++$ )  
    for each transaction  $t$  in transaction table  $T$   
        increment the count of all candidates in  $C_k$  that are contained in  $t$   
     $L_k =$  candidates in  $C_k$  with min_support (frequent)  
     $C_{k+1} =$  candidates generated from  $L_k$ ;  
return  $\cup_k L_k$ ;
```

Special self-join!

The Apriori Algorithm Example (1)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)

Scan T →

C_1	
itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	1
{5}	3

→ L_1

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

↺

C_2	
itemset	
{1 2}	
{1 3}	
{1 5}	
{2 3}	
{2 5}	
{3 5}	

The Apriori Algorithm Example (2)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)

L_2

itemset	sup
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2



C_3

itemset
{2 3 5}


C_2

itemset	sup
{1 2}	1
{1 3}	2
{1 5}	1
{2 3}	2
{2 5}	3
{3 5}	2

Scan T

C_2

itemset
{1 2}
{1 3}
{1 5}
{2 3}
{2 5}
{3 5}



The Apriori Algorithm Example (3)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)



C_3	itemset
	{2 3 5}

Scan T

L_3	itemset	sup
	{2 3 5}	2

The Apriori Algorithm Example (4)

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)

The result of frequent itemsets

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

itemset	sup
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2

itemset	sup
{2 3 5}	2

$$L_1 \cup L_2 \cup L_3$$

Candidates Generation

Advanced

- Suppose the items in L_k are listed in an order
- Step 1: self-joining L_k to get C_{k+1} (In SQL)

INSERT INTO C_{k+1}

SELECT $p.item_1, p.item_2, \dots, p.item_k, q.item_k$

FROM $L_k p, L_k q$

WHERE $p.item_1=q.item_1, \dots, p.item_{k-1}=q.item_{k-1}, p.item_k < q.item_k$

- Step 2: pruning frequent itemsets in C_{k+1}
 - forall *itemsets* c in C_{k+1} do
 - forall *k-subsets* s of c do
 - if (s is not in L_k) then delete c from C_{k+1}

The Previous Example

Advanced

Trans. Table T

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

min_sup=2 (or 50%)

We only need
to match
{2 3} with {2 5}



L_2

itemset	sup
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2


C_3

itemset
{2 3 5}

Example of Candidates Generation

Advanced

- $L_3 = \{abc, abd, acd, ace, bcd\}$
- Self-joining: $L_3 \bowtie L_3$
 - $abcd$ from abc and abd
 - $acde$ from acd and ace
 - No need to match other pairs
- Pruning:
 - $acde$ is removed because ade is not in L_3
- $C_4 = \{abcd\}$
 - Scanning transaction table T is still needed to get the frequencies for items in C_4 (to decide the correct L_4)



Agenda

- Problem definition
- Apriori principle, Apriori algorithm
- Deriving association rules from frequent itemsets


Generating Association Rules from Frequent Itemsets

- Assume that we have discovered the frequent itemsets and their support
- How do we generate association rules?
- Frequent itemsets:

{1}	2
{2}	3
{3}	3
{5}	3
{1,3}	2
{2,3}	2
{2,5}	3
{3,5}	2
{2,3,5}	2

$l \rightarrow$

Not a transaction table!



- For each frequent itemset l , find all nonempty subsets s .
- For each s , generate rule $s \Rightarrow l-s$, if $\text{sup}(l)/\text{sup}(s) \geq \text{min_conf}$

Example: $l = \{2,3,5\}$, $\text{min_conf} = 75\%$

$\{2,3\} \Rightarrow \{5\}$ $2/2=100\%$ ✓
 $\{2,5\} \Rightarrow \{3\}$ $2/3=66.6\%$ ✗
 $\{3,5\} \Rightarrow \{2\}$ $2/2=100\%$ ✓

do the rest as an exercise

Example in Jupyter Notebook

- Library **mlxtend**
 - To install the library: `pip install mlxtend` in Anaconda Prompt
 - `from mlxtend.frequent_patterns import apriori`: frequent itemsets
 - `from mlxtend.frequent_patterns import association_rules`: rules
- Real data
 - `store_data.csv` (in Moodle)
 - `(7501, 20)`
 - 7501 transactions, each having at most 20 items
- `Lecture9_Apriori_mlxtend_stored_data.ipynb`



Performance Bottlenecks of Apriori

Advanced

- Is Apriori fast enough?
- The core of the Apriori algorithm:
 - Use frequent k -itemsets to generate **candidate** frequent $(k+1)$ -itemsets
 - Use full table scan and pattern matching to collect counts for the candidate itemsets
- The bottleneck of Apriori: **candidate generation**
 - Huge candidate sets:
 - A 10^4 -sized frequent 1-itemset will generate 10^7 candidate 2-itemsets
 - To discover a frequent pattern of size 100, e.g., $\{a_1, a_2, \dots, a_{100}\}$, one needs to generate $2^{100} \approx 10^{30}$ candidates.
 - Multiple scans of database table:
 - Needs $(n + 1)$ scans, n is the length of the longest pattern

Methods to Improve Apriori's Efficiency

Advanced

- Transaction reduction
 - A transaction that does not contain any frequent k-itemset is useless in subsequent scans
- Partitioning
 - Any itemset that is potentially frequent in transaction table T must be frequent in at least one of the partitions of T.

Summary

- Association rule definition
 - Support, confidence, lift and association rule
 - Frequent itemsets
 - Steps for association rule mining
- Apriori algorithm
- Deriving association rules from frequent itemsets

Exercises

1. Refer to the transaction table to the right. Say $\text{sup}(ab)=100$

- Determine the possible values of $\text{sup}(a)$

- Conclusion: $\text{sup}(a)$ 100

- *Hint*: Is it possible that $\text{sup}(a)=70$? Why?

- Determine the possible values of $\text{sup}(abc)$

- Conclusion: $\text{sup}(abc)$ 100

- *Hint*: Is it possible that $\text{sup}(abc)=120$? Why?

Choose either “ \leq ” or “ \geq ”

Transaction table
(1000 rows)

TID	Items
1	a,b,c
2	a,c
3	b,e,f
...

2. Write a Jupyter Notebook to find the association rules from the Bread Basket dataset (in Moodle)

- Use the template provided in Moodle (Lecture9_Exercise_BreadBasket_template.ipynb)

Readings and References

- Mandatory readings
 - Association Rule: <https://www.geeksforgeeks.org/association-rule/?ref=lbp>
 - Frequent Itemsets: <https://www.geeksforgeeks.org/frequent-item-set-in-data-set-association-rule-mining/?ref=lbp>
 - Apriori Algorithm: <https://www.geeksforgeeks.org/apriori-algorithm/?ref=lbp>
- Further readings
 - Documentation of mlxtend's frequent
 - http://rasbt.github.io/mlxtend/user_guide/frequent_patterns/apriori/
 - http://rasbt.github.io/mlxtend/user_guide/frequent_patterns/association_rules/
 - Tutorials
 - <https://www.geeksforgeeks.org/implementing-apriori-algorithm-in-python/>
 - <https://www.kaggle.com/code/annettecatherinepaul/apriori-algorithm-association-rule-mining>
 - <https://towardsdatascience.com/understand-and-build-fp-growth-algorithm-in-python-d8b989bab342> (FP-Growth, **advanced**)

Even Further Readings

Advanced

- Readings (**optional. Only if you're interested in theory**)
 - Jiawei Han, Micheline Kamber and Jian Pei. Data Mining: Concepts and Techniques (3rd edition), Elsevier Science Ltd, 2011.
 - Chapters 6 and 7 in the textbook
 - Rakesh Agrawal, Ramakrishnan Srikant: Fast Algorithms for Mining Association Rules in Large Databases. VLDB 1994: 487-499
 - Jiawei Han, Jian Pei, Yiwen Yin: Mining Frequent Patterns without Candidate Generation. SIGMOD 2000: 1-12
- Acknowledgment: Slides are from
 - Margaret H. Dunham (Data Mining: Introductory and Advanced Topics, Prentice Hall, 2002)
 - The HKP textbook
 - Man Lung Yiu and Panagiotis Karras